



UC Berkeley
Teaching Professor
Dan Garcia

CS61C

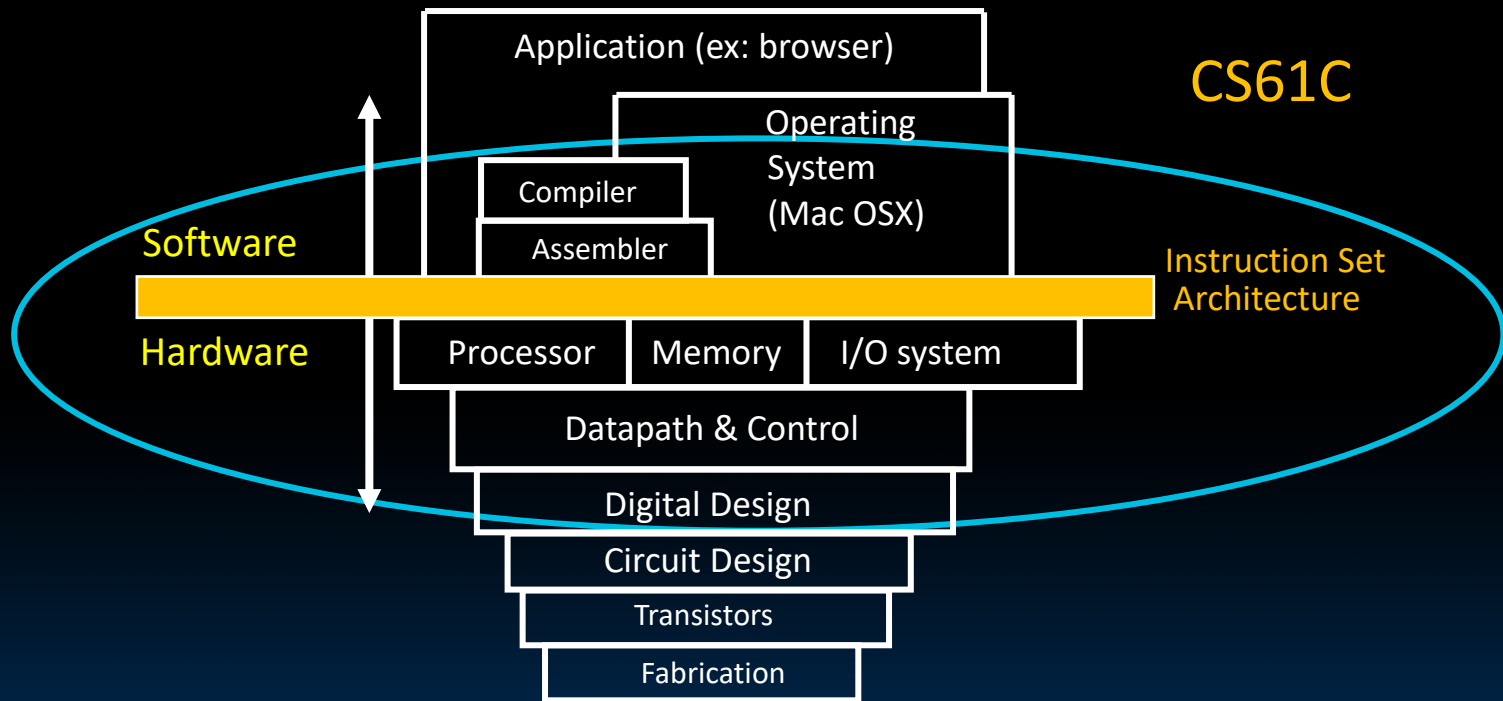
Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley
Professor
Bora Nikolić

RISC-V Processor Design

Machine Structures



New-School Machine Structures

Software

Parallel Requests

Assigned to computer
e.g., Search “Cats”

Parallel Threads

Assigned to core e.g., Lookup, Ads

Parallel Instructions

>1 instruction @ one time
e.g., 5 pipelined instructions

Parallel Data

>1 data item @ one time
e.g., Add of 4 pairs of words

Hardware descriptions

All gates work in parallel at same time

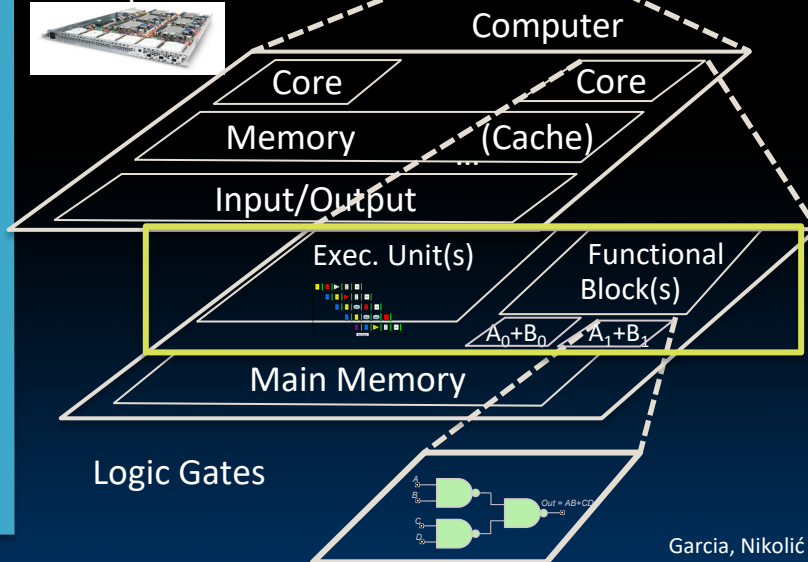
Harness
Parallelism &
Achieve High
Performance

Hardware

Warehouse
Scale
Computer



Smart
Phone



Great Idea #1: Abstraction (Levels of Representation/Interpretation)

High Level Language
Program (e.g., C)

```
temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;
```

Compiler

Assembly Language
Program (e.g., RISC-V)

```
lw    x3, 0(x10)
lw    x4, 4(x10)
sw    x4, 0(x10)
sw    x3, 4(x10)
```

Assembler

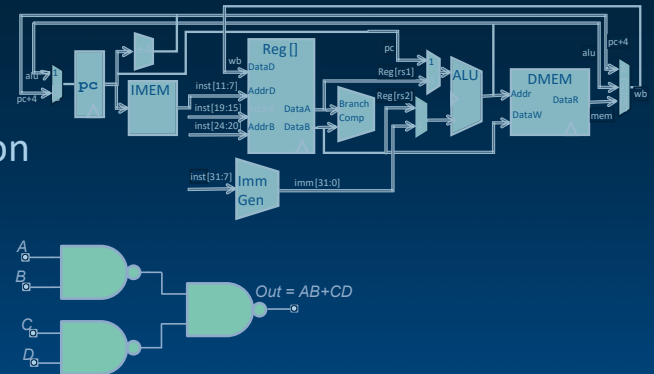
Machine Language
Program (RISC-V)

```
1000 1101 1110 0010 0000 0000 0000 0000
1000 1110 0001 0000 0000 0000 0000 0100
1010 1110 0001 0010 0000 0000 0000 0000
1010 1101 1110 0010 0000 0000 0000 0100
```

Hardware Architecture Description
(e.g., block diagrams)

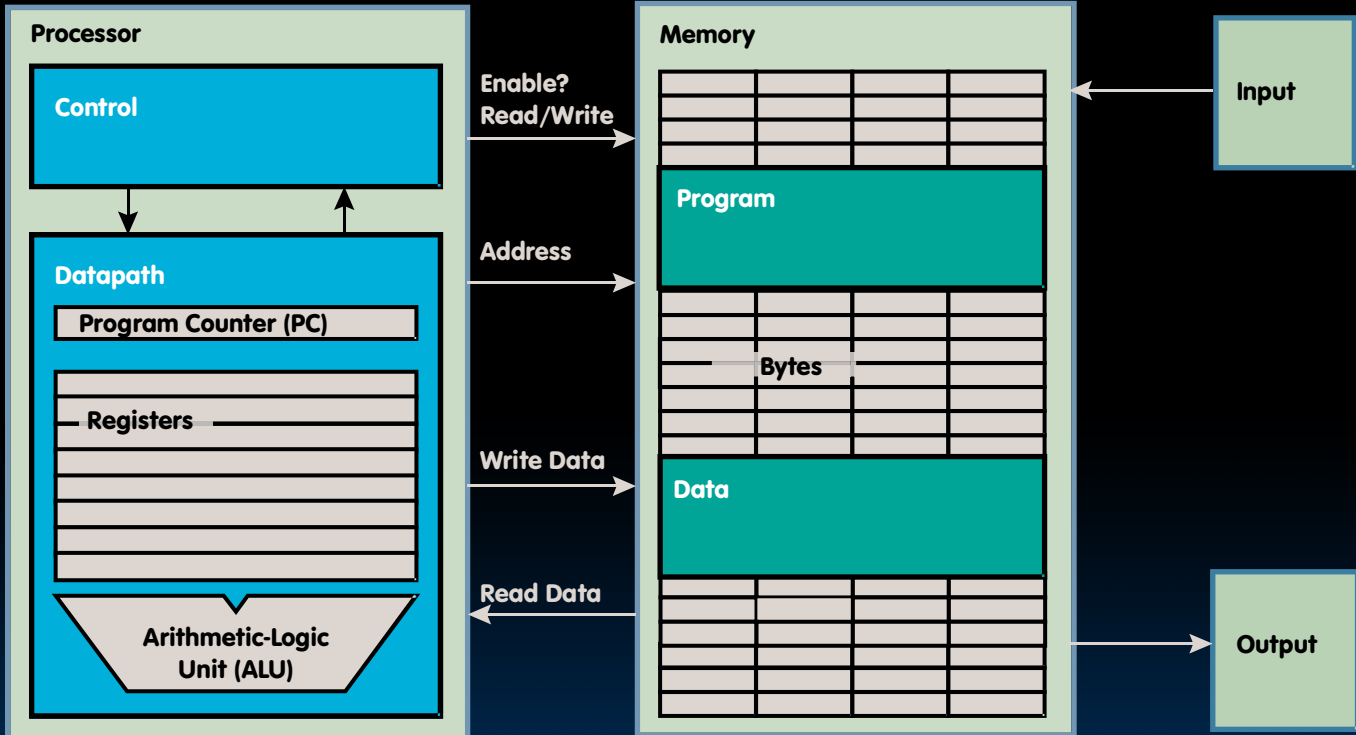
Architecture Implementation

Logic Circuit Description
(Circuit Schematic Diagrams)





Our Single-Core Processor So Far...







The CPU

- **Processor (CPU):** the active part of the computer that does all the work (data manipulation and decision-making)
- **Datapath:** portion of the processor that contains hardware necessary to perform operations required by the processor (the brawn)
- **Control:** portion of the processor (also in hardware) that tells the datapath what needs to be done (the brain)

Need to Implement All RV32I Instructions

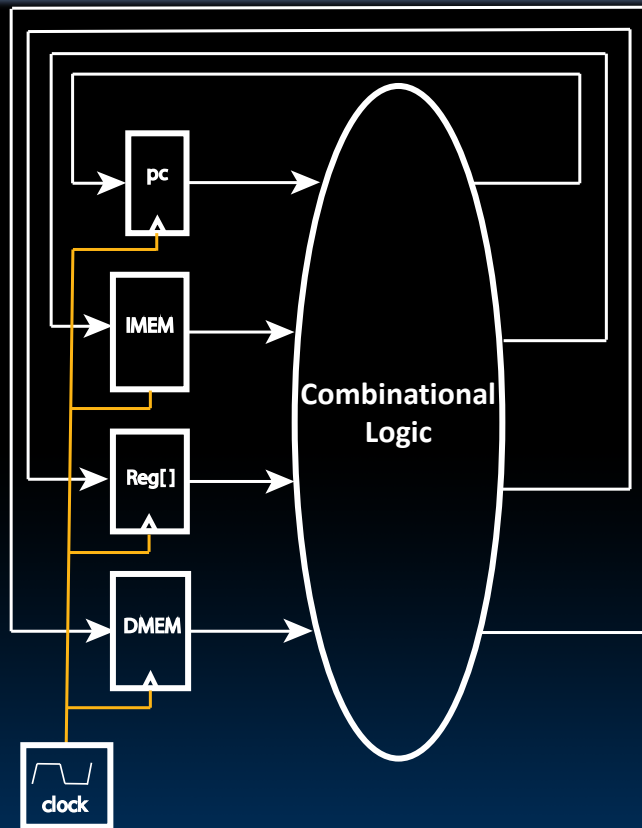
Open  RISC-V Reference Card							
Base Integer Instructions: RV32I							
Category	Name	Fmt	RV32I Base	Category	Name	Fmt	RV32I Base
Shifts	Shift Left Logical	R	SLL rd,rs1,rs2	Loads	Load Byte	I	LB rd,rs1,imm
	Shift Left Log. Imm.	I	SLLI rd,rs1,shamt		Load Halfword	I	LH rd,rs1,imm
	Shift Right Logical	R	SRL rd,rs1,rs2		Load Byte Unsigned	I	LBU rd,rs1,imm
	Shift Right Log. Imm.	I	SRLI rd,rs1,shamt		Load Half Unsigned	I	LHU rd,rs1,imm
	Shift Right Arithmetic	R	SRA rd,rs1,rs2		Load Word	I	LW rd,rs1,imm
	Shift Right Arith. Imm.	I	SRAI rd,rs1,shamt	Stores	Store Byte	S	SB rs1,rs2,imm
Arithmetic	ADD	R	ADD rd,rs1,rs2		Store Halfword	S	SH rs1,rs2,imm
	ADD Immediate	I	ADDI rd,rs1,imm		Store Word	S	SW rs1,rs2,imm
	SUBtract	R	SUB rd,rs1,rs2	Branches	Branch =	B	BEQ rs1,rs2,imm
	Load Upper Imm	U	LUI rd,imm		Branch ≠	B	BNE rs1,rs2,imm
	Add Upper Imm to PC	U	AUIPC rd,imm		Branch <	B	BLT rs1,rs2,imm
Logical	XOR	R	XOR rd,rs1,rs2		Branch ≥	B	BGE rs1,rs2,imm
	XOR Immediate	I	XORI rd,rs1,imm		Branch < Unsigned	B	BLTU rs1,rs2,imm
	OR	R	OR rd,rs1,rs2		Branch ≥ Unsigned	B	BGEU rs1,rs2,imm
	OR Immediate	I	ORI rd,rs1,imm	Jump & Link	J&L	J	JAL rd,imm
	AND	R	AND rd,rs1,rs2		Jump & Link Register	I	JALR rd,rs1,imm
Compare	AND Immediate	I	ANDI rd,rs1,imm	Synch	Synch thread	I	FENCE
	Set <	R	SLT rd,rs1,rs2				
	Set < Immediate	I	SLTI rd,rs1,imm	Environment	CALL	I	ECALL
	Set < Unsigned	R	SLTU rd,rs1,rs2		BREAK	I	EBREAK
	Set < Imm Unsigned	I	SLTIU rd,rs1,imm				

Not in
61C



Building a RISC-V Processor

One-Instruction-Per-Cycle RISC-V Machine



- On every tick of the clock, the computer executes one instruction
- Current state outputs drive the inputs to the combinational logic, whose outputs settle at the values of the state before the next clock edge
- At the rising clock edge, all the state elements are updated with the combinational logic outputs, and execution moves to the next clock cycle



Stages of the Datapath : Overview

- Problem: a single, “monolithic” block that “executes an instruction” (performs all necessary operations beginning with fetching the instruction) would be too bulky and inefficient
- Solution: break up the process of “executing an instruction” into stages, and then connect the stages to create the whole datapath
 - smaller stages are easier to design
 - easy to optimize (change) one stage without touching the others (modularity)

Five Stages of the Datapath


- Stage 1: *Instruction Fetch (IF)*
- Stage 2: *Instruction Decode (ID)*
- Stage 3: *Execute (EX) - ALU (Arithmetic-Logic Unit)*
- Stage 4: *Memory Access (MEM)*
- Stage 5: *Write Back to Register (WB)*

Basic Phases of Instruction Execution



Datapath Components: Combinational


- Combinational elements



- Storage elements + clocking methodology
- Building blocks

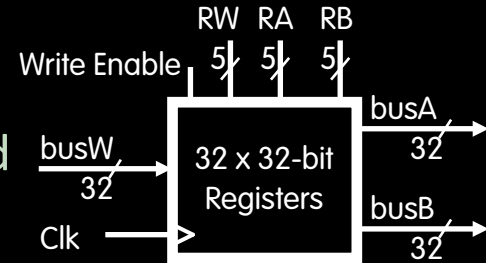
Datapath Elements: State and Sequencing (1/3)

- Register
- Write Enable:
 - Low (or deasserted) (0): Data Out will not change
 - Asserted (1): Data Out will become Data In on positive edge of clock




Datapath Elements: State and Sequencing (2/3)

- Register file (regfile, RF) consists of 32 registers:
 - Two 32-bit output busses: busA and busB
 - One 32-bit input bus: busW
- Register is selected by:
 - RA (number) selects the register to put on busA (data)
 - RB (number) selects the register to put on busB (data)
 - RW (number) selects the register to be written via busW (data) when Write Enable is 1
- Clock input (Clk)
 - Clk input is a factor ONLY during write operation
 - During read operation, behaves as a combinational logic block:
 - RA or RB valid \Rightarrow busA or busB valid after “access time.”



Datapath Elements: State and Sequencing (3/3)

- “Magic” Memory
 - One input bus: Data In
 - One output bus: Data Out
- Memory word is found by:
 - For Read: Address selects the word to put on Data Out
 - For Write: Set Write Enable = 1: address selects the memory word to be written via the Data In bus
- Clock input (CLK)
 - CLK input is a factor ONLY during write operation
 - During read operation, behaves as a combinational logic block: Address valid \Rightarrow Data Out valid after “access time”



State Required by RV32I ISA (1/2)

Each instruction during execution reads and updates the state of : (1) Registers, (2) Program counter, (3) Memory

- Registers (**x0** . . **x31**)
 - Register file (*regfile*) **Reg** holds 32 registers x 32 bits/register: **Reg[0]** . . **Reg[31]**
 - First register read specified by **rs1** field in instruction
 - Second register read specified by **rs2** field in instruction
 - Write register (destination) specified by **rd** field in instruction
 - **x0** is always 0 (writes to **Reg[0]** are ignored)
- Program Counter (**PC**)
 - Holds address of current instruction

State Required by RV32I ISA (2/2)

- Memory (**MEM**)
 - Holds both instructions & data, in one 32-bit byte-addressed memory space
 - We'll use separate memories for instructions (**IMEM**) and data (**DMEM**)
 - *These are placeholders for instruction and data caches*
 - Instructions are read (*fetch*) from instruction memory (assume **IMEM** read-only)
 - Load/store instructions access data memory

R-Type Add Datapath

Review: R-Type Instructions

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	10	9	8	7	6	5	4	3	2	1	0
R-format : ALU																														
[31:25]						[24:20]						[19:15]					[14:12]				[11:7]					[6:0]				
7						5						5					3				5					7				
func7						rs2						rs1					func3				rd					opcode				
0000000						rs2						rs1					000 : ADD				rd					0110011:OP-R				
0100000						rs2						rs1					000 : SUB				rd					0110011:OP-R				
0000000						rs2						rs1					001 : SLL				rd					0110011:OP-R				
0000000						rs2						rs1					010 : SLT				rd					0110011:OP-R				
0000000						rs2						rs1					011 : SLTU				rd					0110011:OP-R				
0000000						rs2						rs1					100 : XOR				rd					0110011:OP-R				
0000000						rs2						rs1					101 : SRL				rd					0110011:OP-R				
0100000						rs2						rs1					101 : SRA				rd					0110011:OP-R				
0000000						rs2						rs1					110 : OR				rd					0110011:OP-R				
0000000						rs2						rs1					111 : AND				rd					0110011:OP-R				

- E.g. Addition/subtraction **add rd, rs1, rs2**

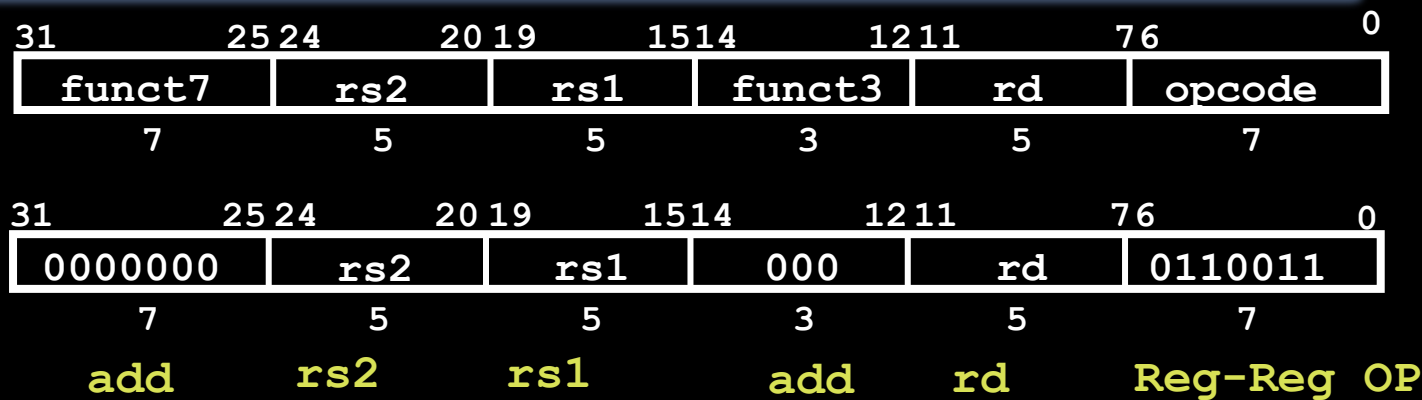
$$R[rd] = R[rs1] + R[rs2]$$

sub rd, rs1, rs2

$$R[rd] = R[rs1] - R[rs2]$$



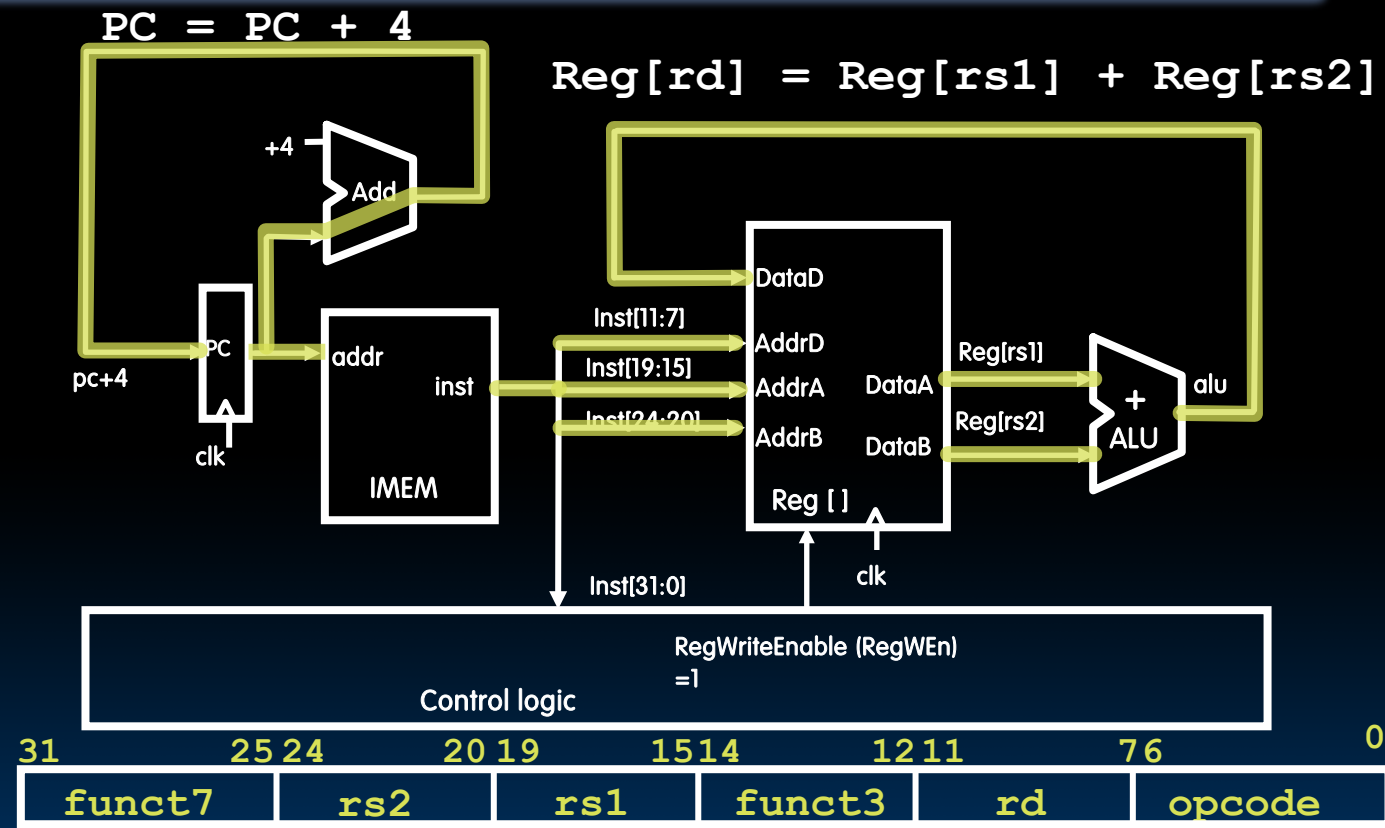
Implementing the add instruction



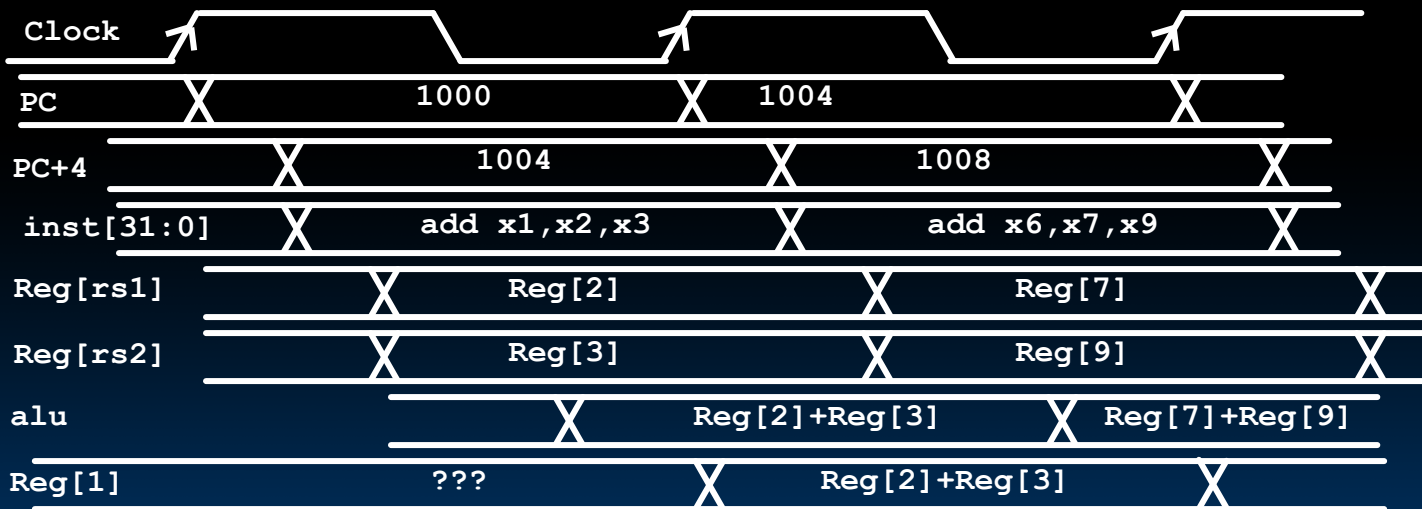
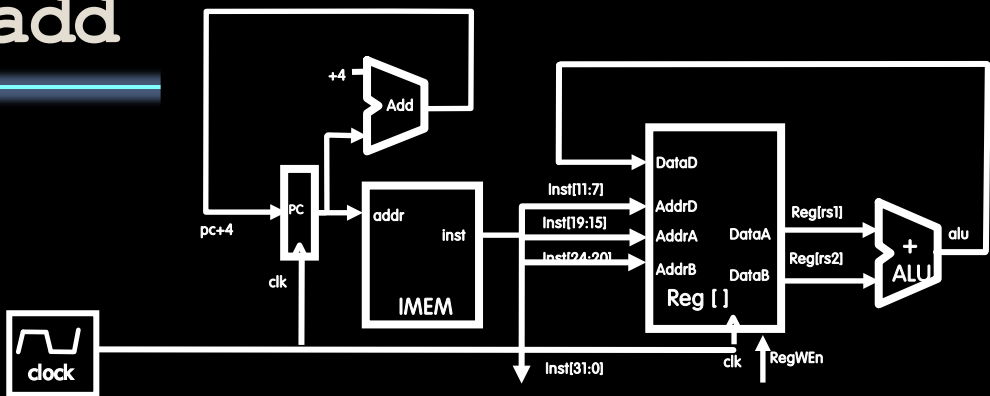
add rd, rs1, rs2

- Instruction makes two changes to machine's state:
 - $\text{Reg}[\text{rd}] = \text{Reg}[\text{rs1}] + \text{Reg}[\text{rs2}]$
 - $\text{PC} = \text{PC} + 4$

Datapath for add



Timing Diagram for add



Sub Datapath

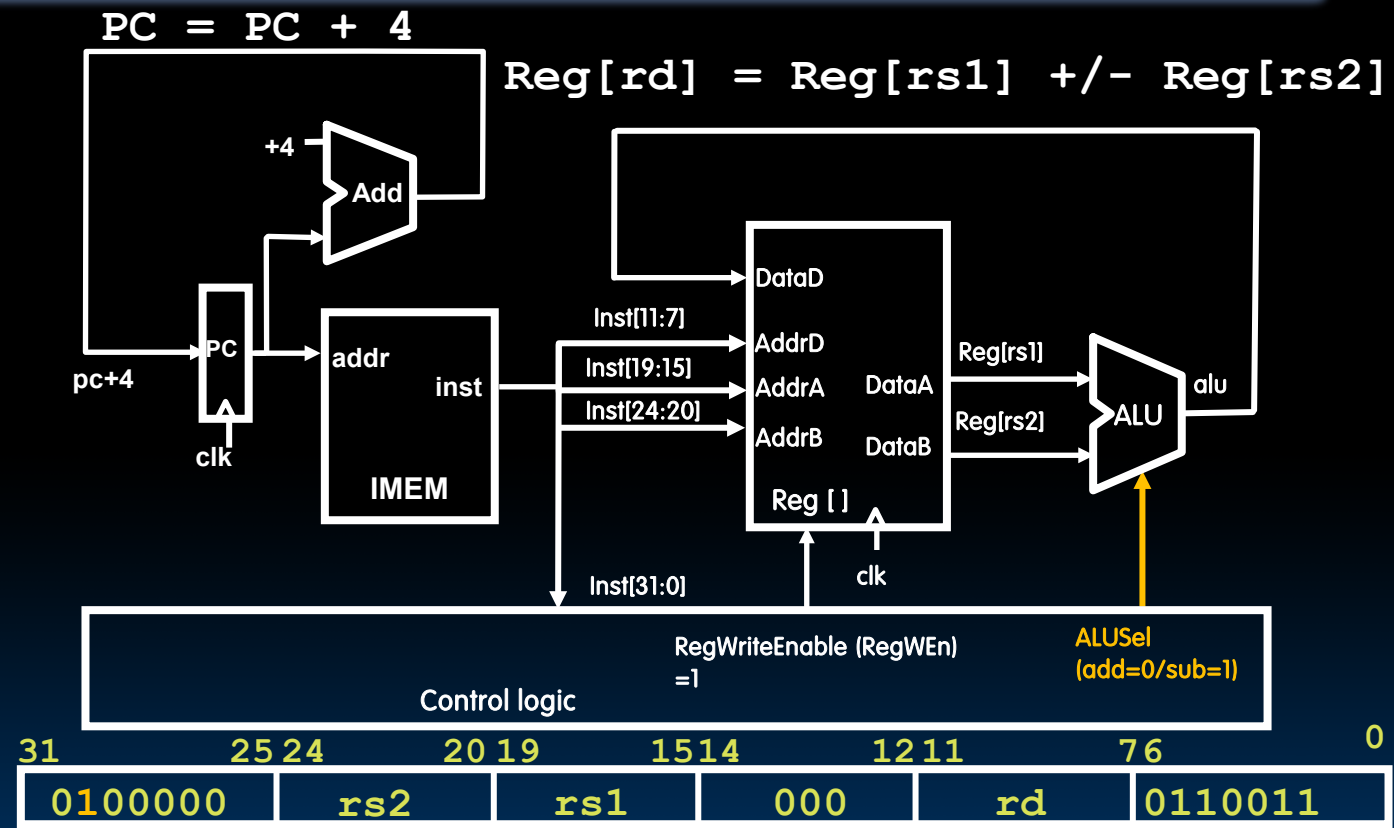
Implementing the sub instruction

0000000	rs2	rs1	000	rd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub

sub rd, rs1, rs2

- Almost the same as add, except now have to subtract operands instead of adding them
- **inst[30]** selects between add and subtract

Datapath for add/sub



Implementing Other R-Format Instructions

0000000	rs2	rs1	000	rd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub
0000000	rs2	rs1	001	rd	0110011	sll
0000000	rs2	rs1	010	rd	0110011	slt
0000000	rs2	rs1	011	rd	0110011	sltu
0000000	rs2	rs1	100	rd	0110011	xor
0000000	rs2	rs1	101	rd	0110011	srl
0100000	rs2	rs1	101	rd	0110011	sra
0000000	rs2	rs1	110	rd	0110011	or
0000000	rs2	rs1	111	rd	0110011	and


All implemented by decoding funct3 and funct7 fields and selecting appropriate ALU function

Datapath With Immediates

Implementing I-Format - addi instruction

- RISC-V Assembly Instruction:

addi x15, x1, -50



imm=-50

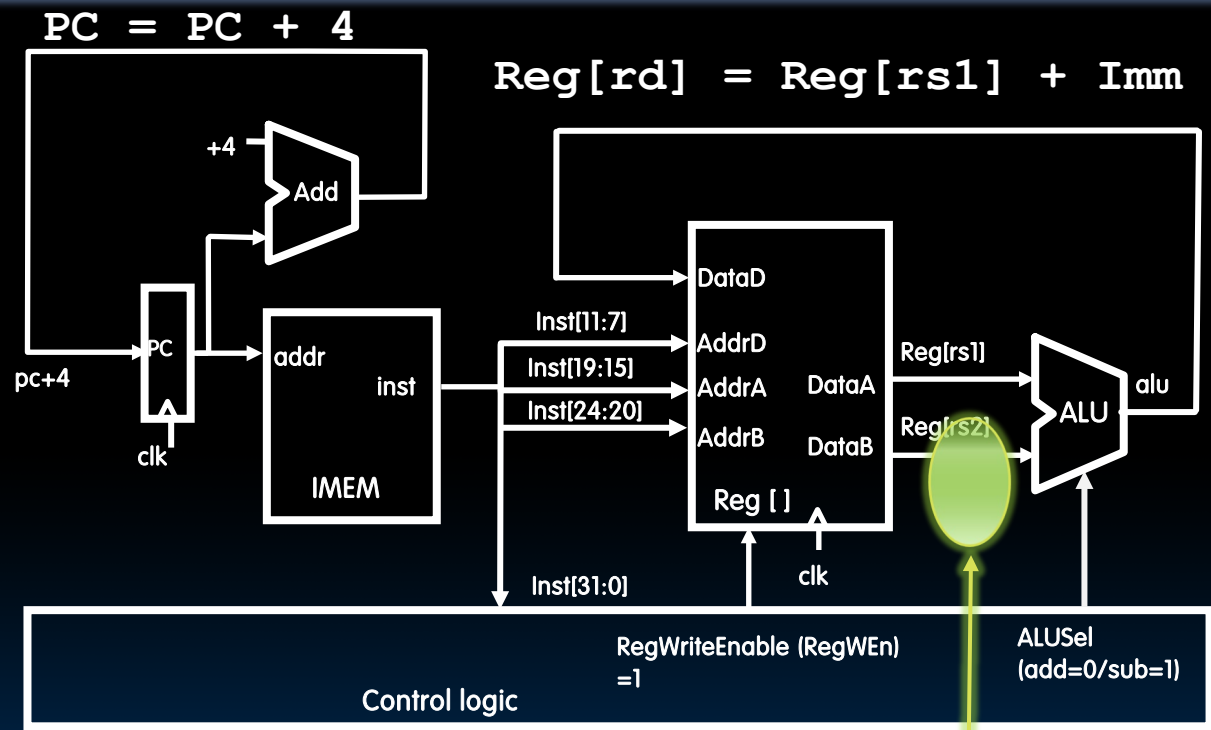
rs1=1

add

rd=15

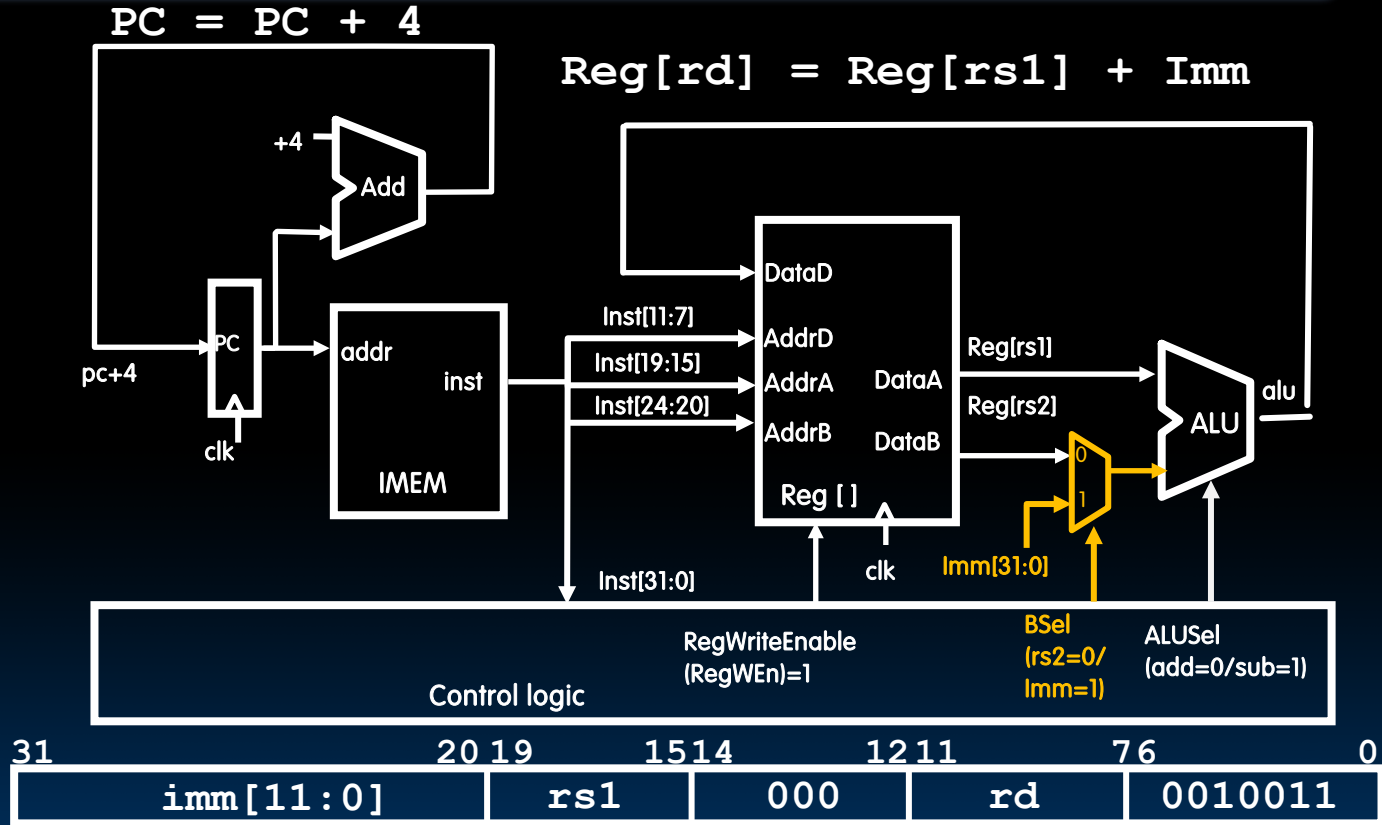
OP-Imm

Datapath for add/sub

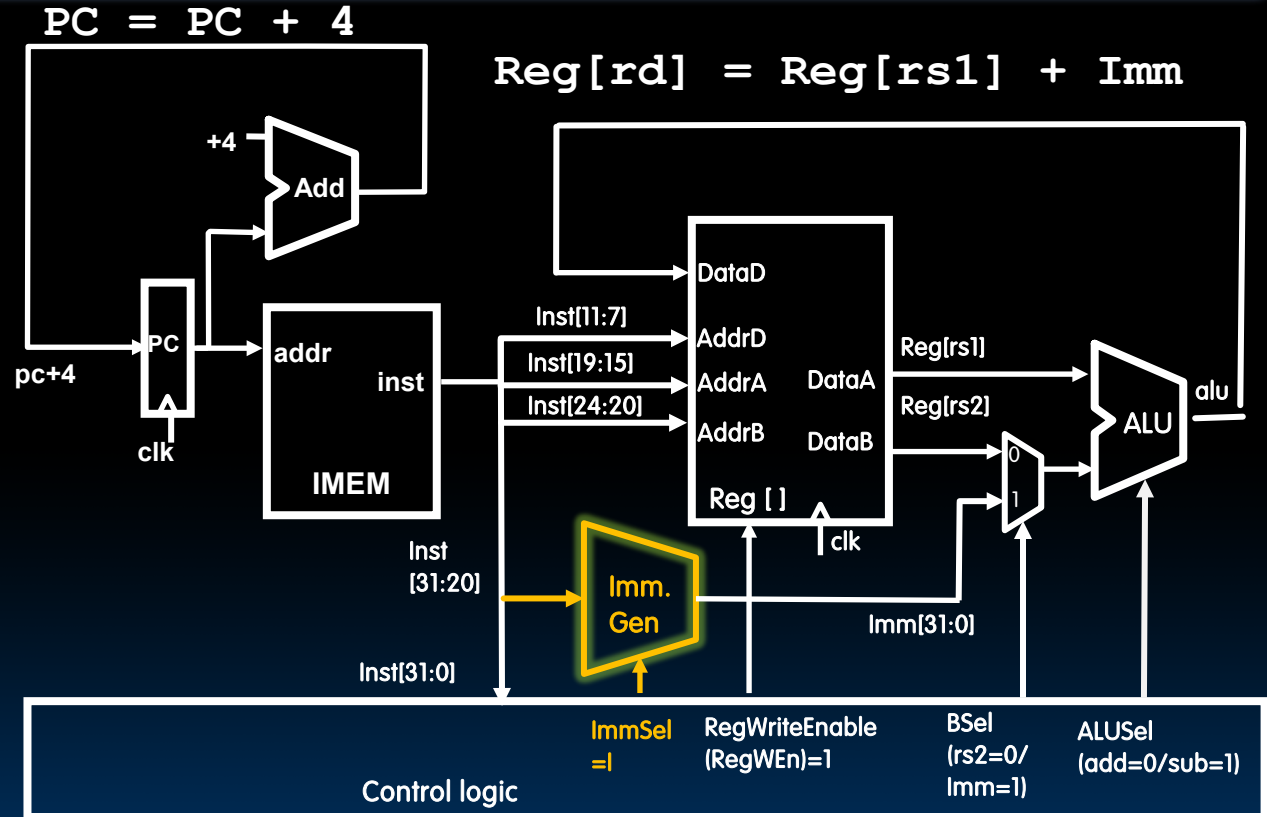


Immediate should
be here

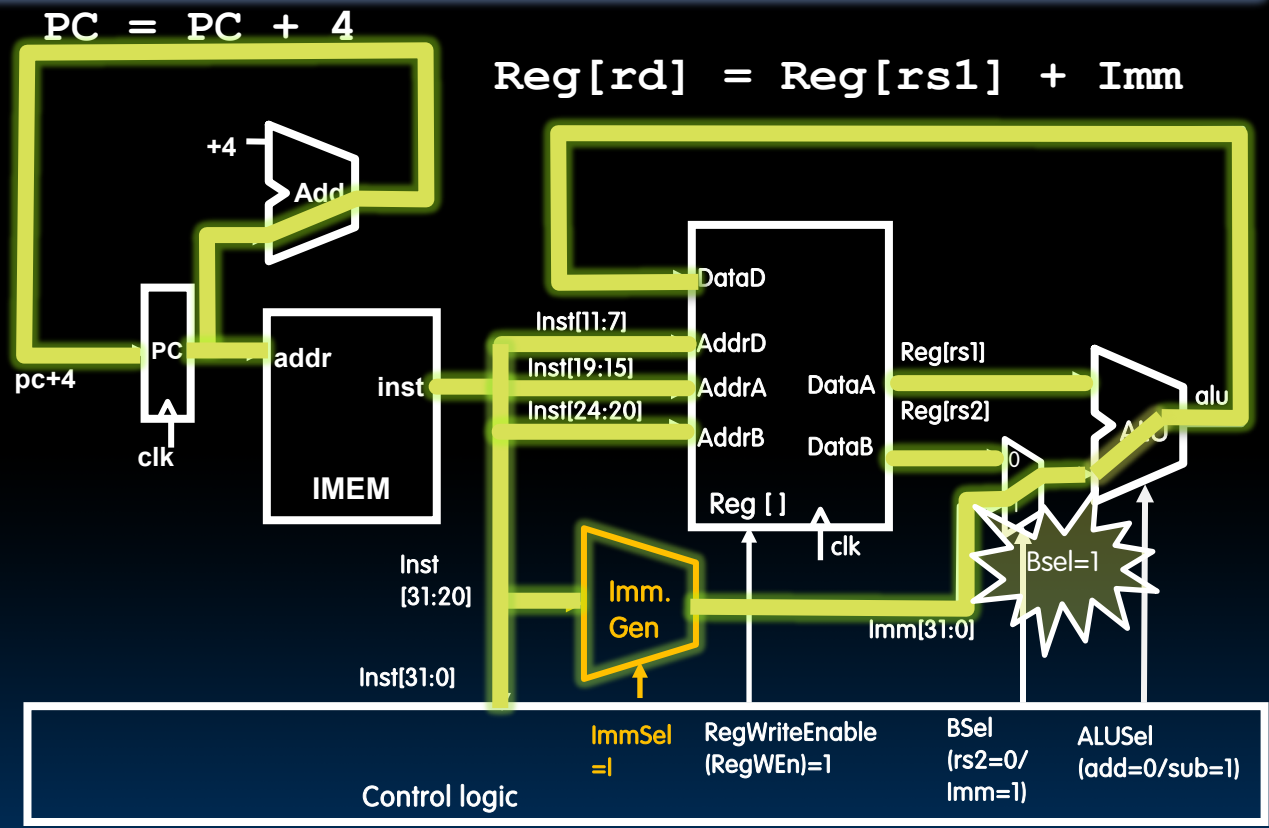
Adding addi to Datapath



Adding addi to Datapath




Adding addi to Datapath



I-Format Immediates

-inst[31]-



- High 12 bits of instruction (**inst[31:20]**) copied to low 12 bits of immediate (**imm[11:0]**)
- Immediate is sign-extended by copying value of **inst[31]** to fill the upper 20 bits of the immediate value (**imm[31:12]**)

Adding addi to Datapath

