

Hierarchical Page Tables


Size of Page Tables

- E.g., 32-Bit virtual address, 4-KiB pages
 - Single page table size:
 - 4×2^{20} Bytes = 4-MiB
 - 0.1% of 4-GiB memory
 - Total size for 256 processes (each needs a page table)
 - $256 \times 4 \times 2^{20}$ Bytes = $256 \times 4\text{-MiB} = 1\text{-GiB}$
 - 25% of 4-GiB memory!
- What about 64-bit addresses?

How can we keep the size of page tables “reasonable”?

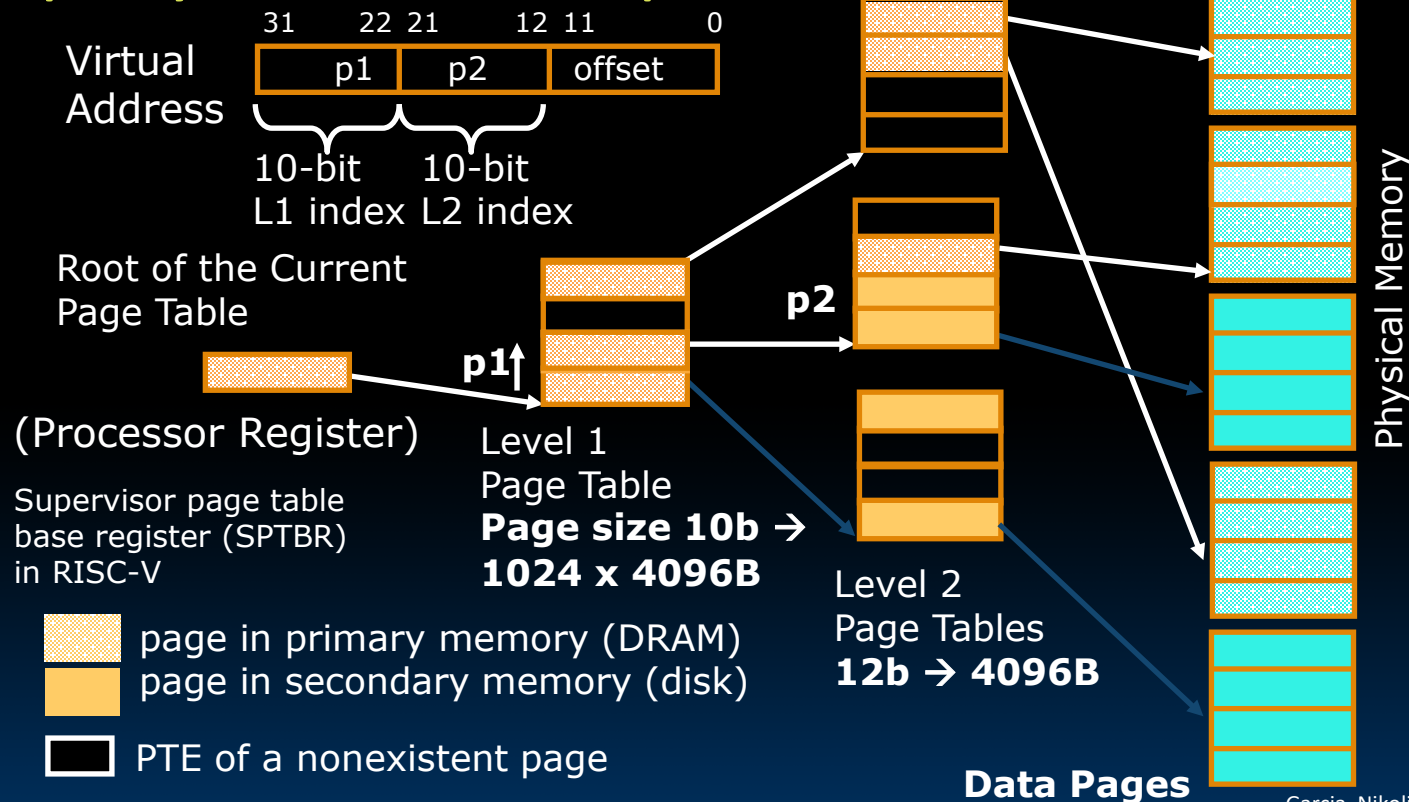
Options for Page Tables

- Increase page size
 - E.g., doubling page size cuts PT size in half
 - At the expense of potentially wasted memory
- Hierarchical page tables
 - With decreasing page size
- Most programs use only fraction of memory
 - Split PT in two (or more) parts
 - This is done in RISC-V




Hierarchical Page Table

Exploits Sparsity of Virtual Address Space Use



Example: 32-b RISC-V



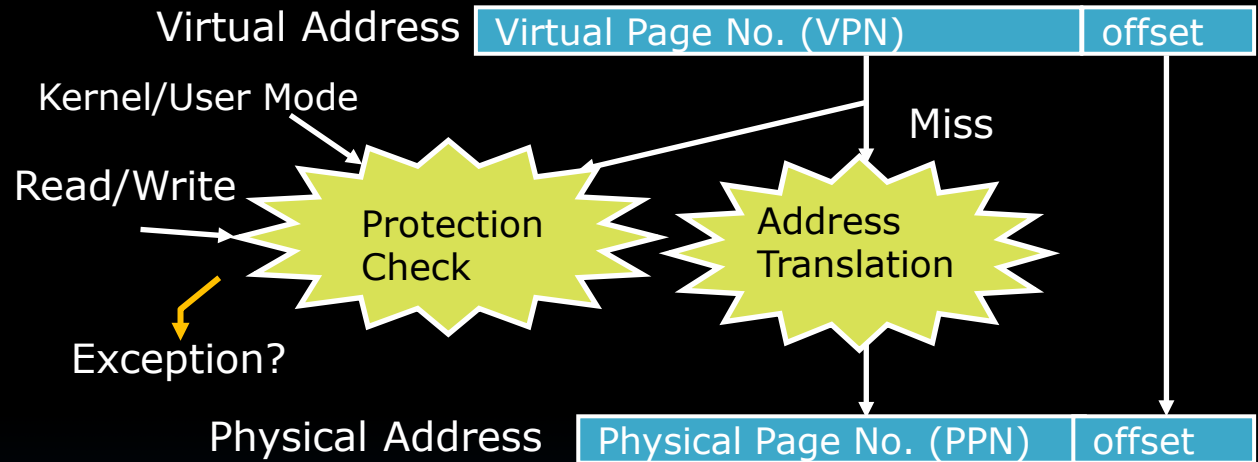
- VPN: Virtual Page Number
- PPN: Physical Page Number
- Page Table Entry (PTE) is 32b and contains:
 - PPN[1], PPN[0]
 - Status bits for protection and usage (read, write, exec), validity, etc.



R= 0, W=0, X = 0 points to next level page table;
otherwise it is a leaf PTE

Translation Lookaside Buffers

Address Translation and Protection



- Every instruction and data access needs address translation and protection checks

Good VM design should be fast (~one cycle) and space efficient

Translation Lookaside Buffers (TLB)

Address translation is very expensive!


In a single-level page table, each reference becomes two memory accesses

In a two-level page table, each reference becomes three memory accesses

Solution: *Cache some translations in TLB*

TLB hit → *Single-Cycle Translation*

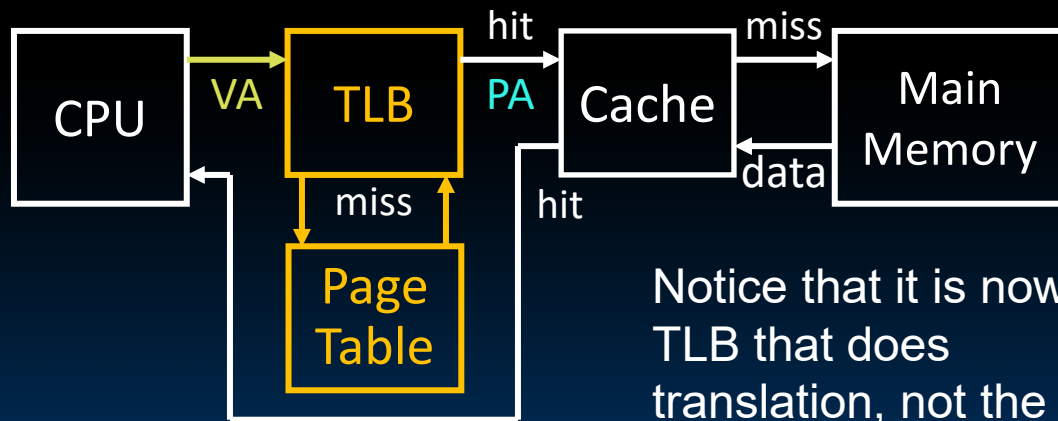
TLB miss → *Page-Table Walk to refill*



- Typically 32-128 entries, usually fully associative
 - Each entry maps a large page, hence less spatial locality across pages → more likely that two entries conflict
 - Sometimes larger TLBs (256-512 entries) are 4-8 way set-associative
 - Larger systems sometimes have multi-level (L1 and L2) TLBs
- Random or FIFO replacement policy
- “TLB Reach”: Size of largest virtual address space that can be simultaneously mapped by TLB

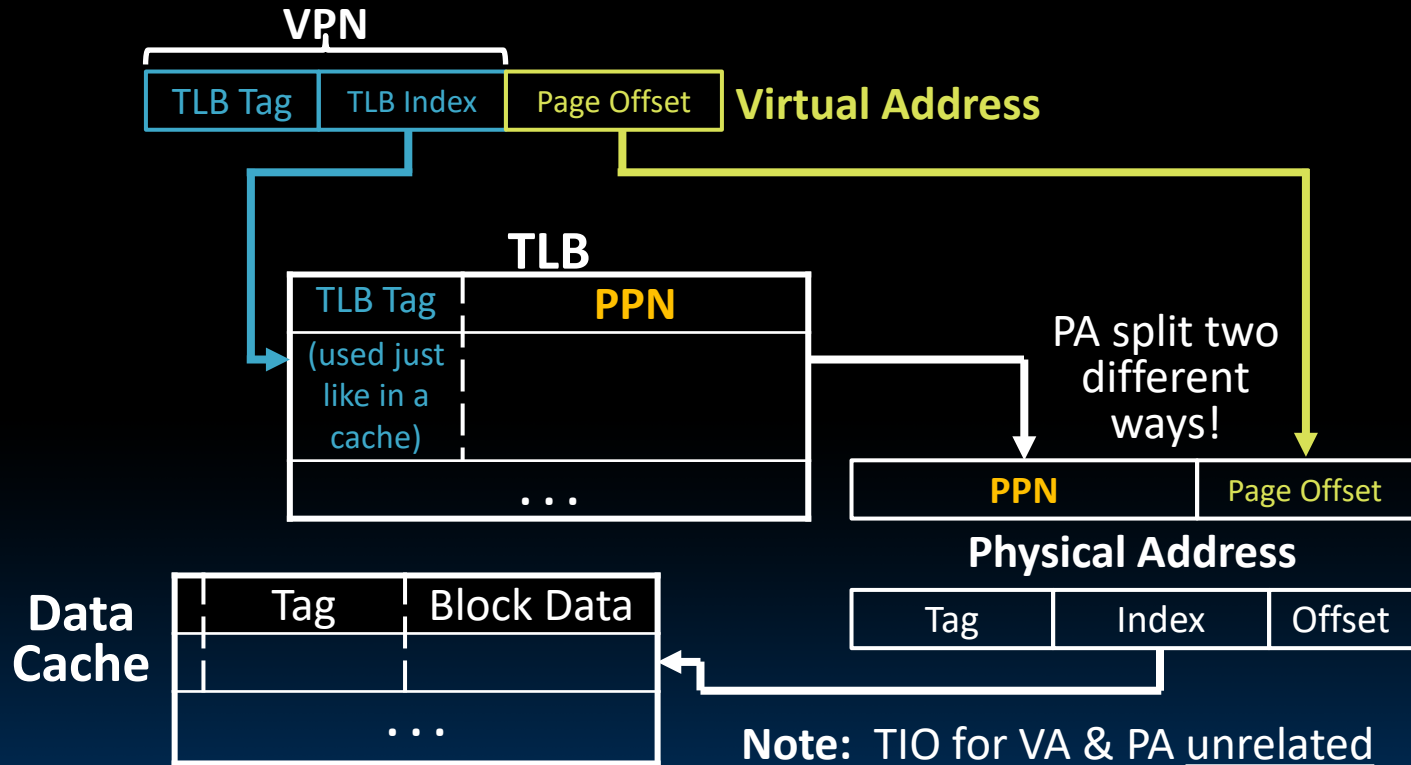
Where Are TLBs Located?

- Which should we check first: Cache or TLB?
 - Can cache hold requested data if corresponding page is not in physical memory? **No**
 - With TLB first, does cache receive VA or PA? **PA**



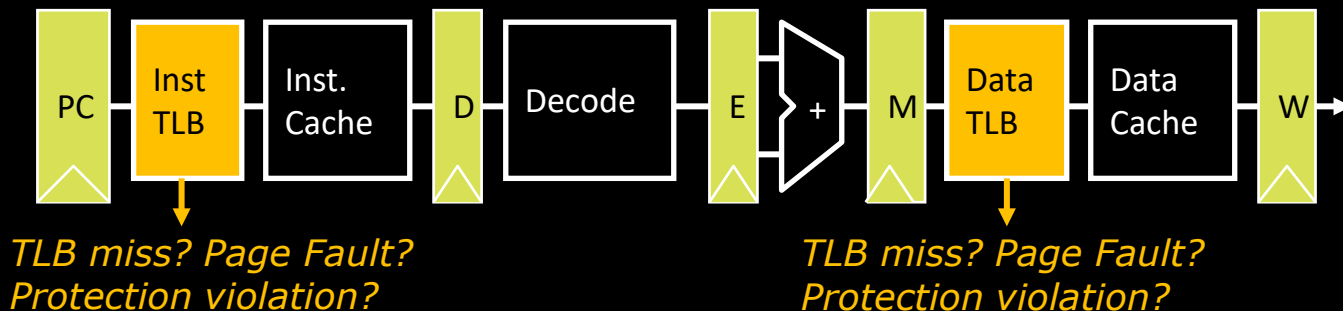
Notice that it is now the TLB that does translation, not the Page Table!

Address Translation Using TLB



TLBs in Datapath

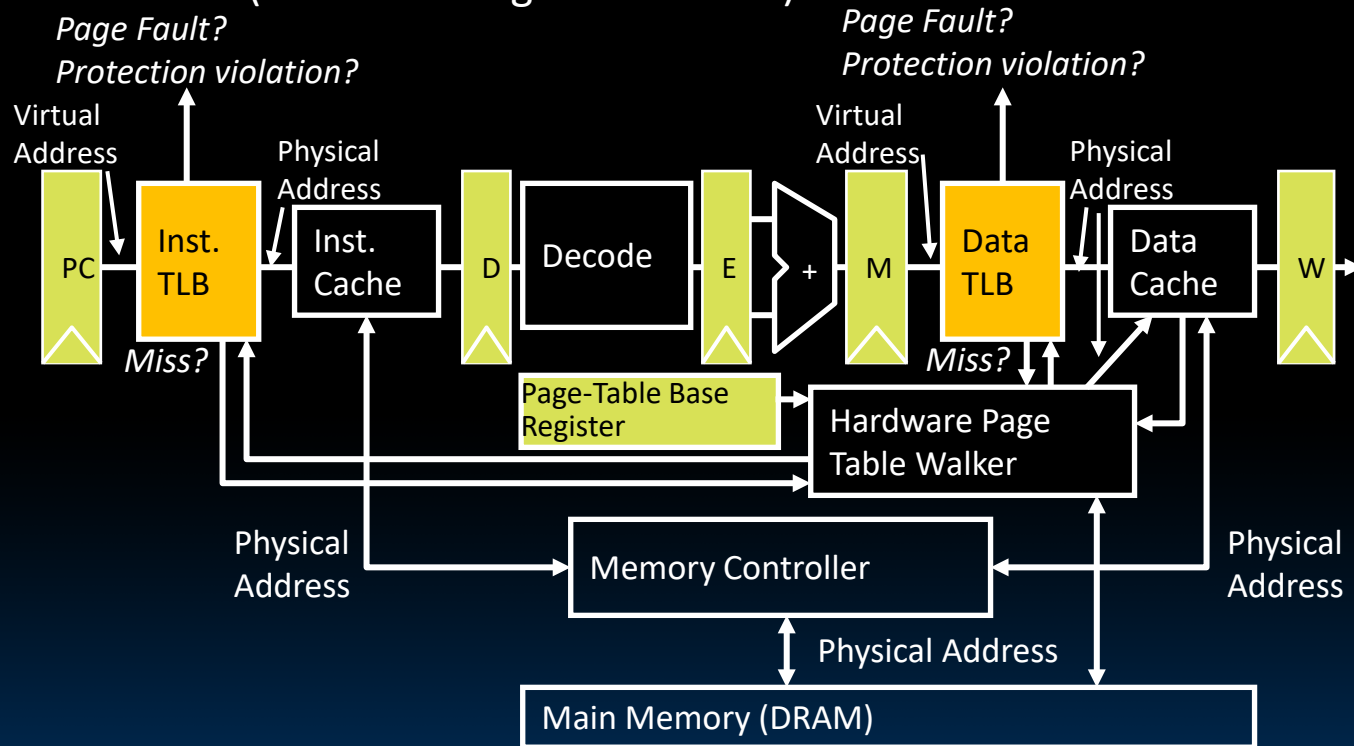
VM-related Events in Pipeline



- Handling a TLB miss needs a hardware or software mechanism to refill TLB
 - Usually done in hardware
- Handling a page fault (e.g., page is on disk) needs a *precise trap* so software handler can easily resume after retrieving page
- Protection violation may abort process

Page-Based Virtual-Memory Machine

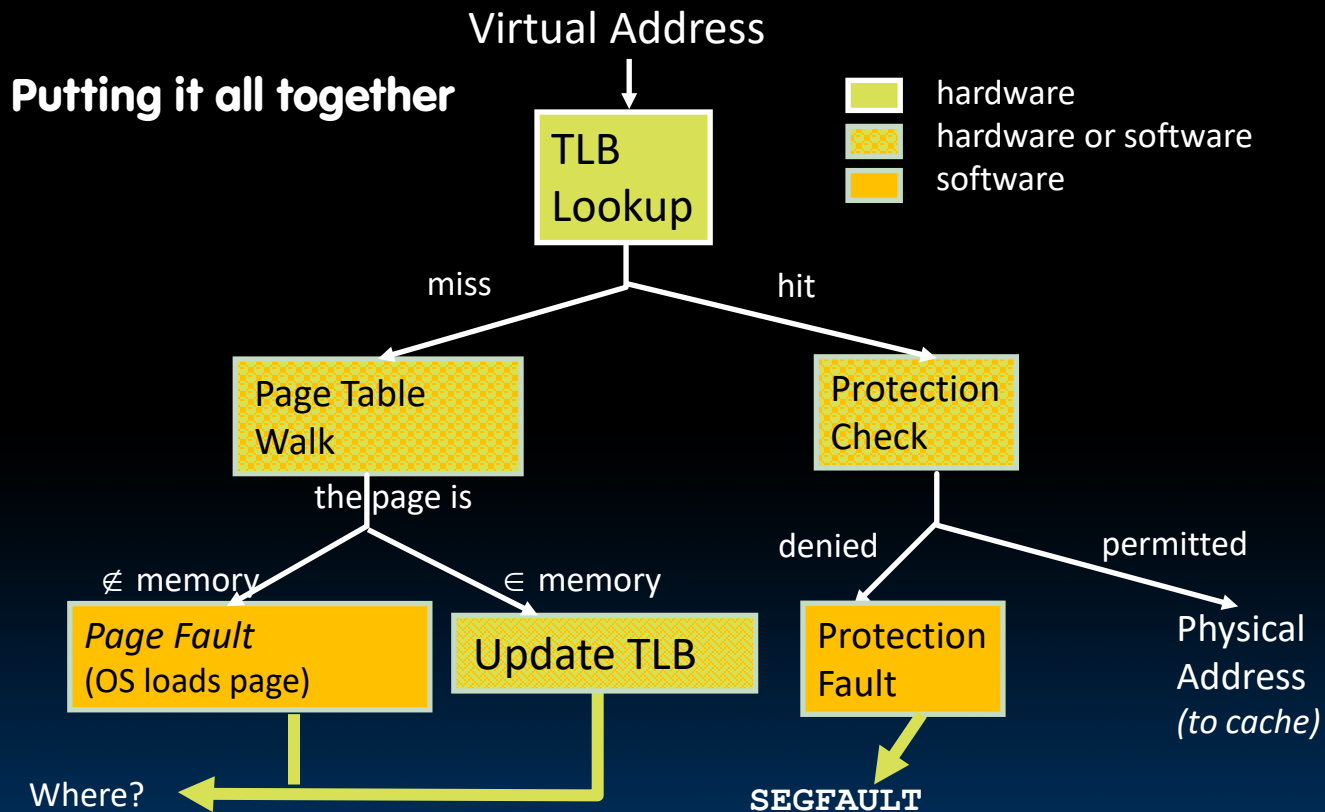
(Hardware Page-Table Walk)



- Assumes page tables held in untranslated physical memory

Address Translation

Putting it all together



Illusion of a large, private, uniform store

Protection & Privacy


Several users/processes, each with their private address space

Demand Paging

Provides the ability to run programs larger than the primary memory

Hides differences in machine configurations

The price is address translation on each memory reference





Review: Context Switching

- How does a single processor run many programs at once?
- **Context switch:** Changing of internal state of processor (switching between processes)
 - Save register values (and PC) and change value in Supervisor Page Table Base register (SPTBR)
- What happens to the TLB?
 - Current entries are for different process
 - Set all entries to invalid on context switch

VM

Performance



Comparing the Cache and VM

Cache version

Block or Line

Miss

Block Size: 32-64B

Placement:

Direct Mapped,
N-way Set Associative

Replacement:

LRU or Random

Write Thru or Back

Virtual Memory version

Page

Page Fault

Page Size: 4K-8KiB

Fully Associative


Least Recently Used
(LRU), FIFO, random

Write Back

VM Performance

- Virtual Memory is the level of the memory hierarchy that sits *below* main memory
 - TLB comes *before* cache, but affects transfer of data from disk to main memory
 - Previously we assumed main memory was lowest level, now we just have to account for disk accesses
- Same CPI, AMAT equations apply, but now treat main memory like a mid-level cache

Typical Performance Stats



Caching

- cache entry
- cache block (≈ 32 -64 bytes)
- cache miss rate (1% to 20%)
- cache hit (≈ 1 cycle)
- cache miss (≈ 100 cycles)

Demand paging

- page frame
- page (≈ 4 Ki bytes)
- page miss rate ($< 0.001\%$)
- page hit (≈ 100 cycles)
- page miss (≈ 5 M cycles)

Impact of Paging on AMAT (1/2)

- Memory Parameters:
 - L1 cache hit = 1 clock cycles, hit 95% of accesses
 - L2 cache hit = 10 clock cycles, hit 60% of L1 misses
 - DRAM = 200 clock cycles (≈ 100 nanoseconds)
 - Disk = 20,000,000 clock cycles (≈ 10 milliseconds)
- Average Memory Access Time (no paging):
 - $1 + 5\% \times 10 + 5\% \times 40\% \times 200 = 5.5$ clock cycles
- Average Memory Access Time (with paging):
 - 5.5 (AMAT with no paging) + ?

Impact of Paging on AMAT (2/2)

- Average Memory Access Time (with paging) =
 - $5.5 + 5\% \times 40\% \times (1 - \text{HR}_{\text{Mem}}) \times 20,000,000$
- AMAT if $\text{HR}_{\text{Mem}} = 99\%$?
 - $5.5 + 0.02 \times 0.01 \times 20,000,000 = 4005.5$ ($\approx 728\times$ slower)
 - 1 in 20,000 memory accesses goes to disk: 10 sec program takes 2 hours!
- AMAT if $\text{HR}_{\text{Mem}} = 99.9\%$?
 - $5.5 + 0.02 \times 0.001 \times 20,000,000 = 405.5$
- AMAT if $\text{HR}_{\text{Mem}} = 99.9999\%$
 - $5.5 + 0.02 \times 0.000001 \times 20,000,000 = 5.9$